

• the next steps Django

Django is great*!

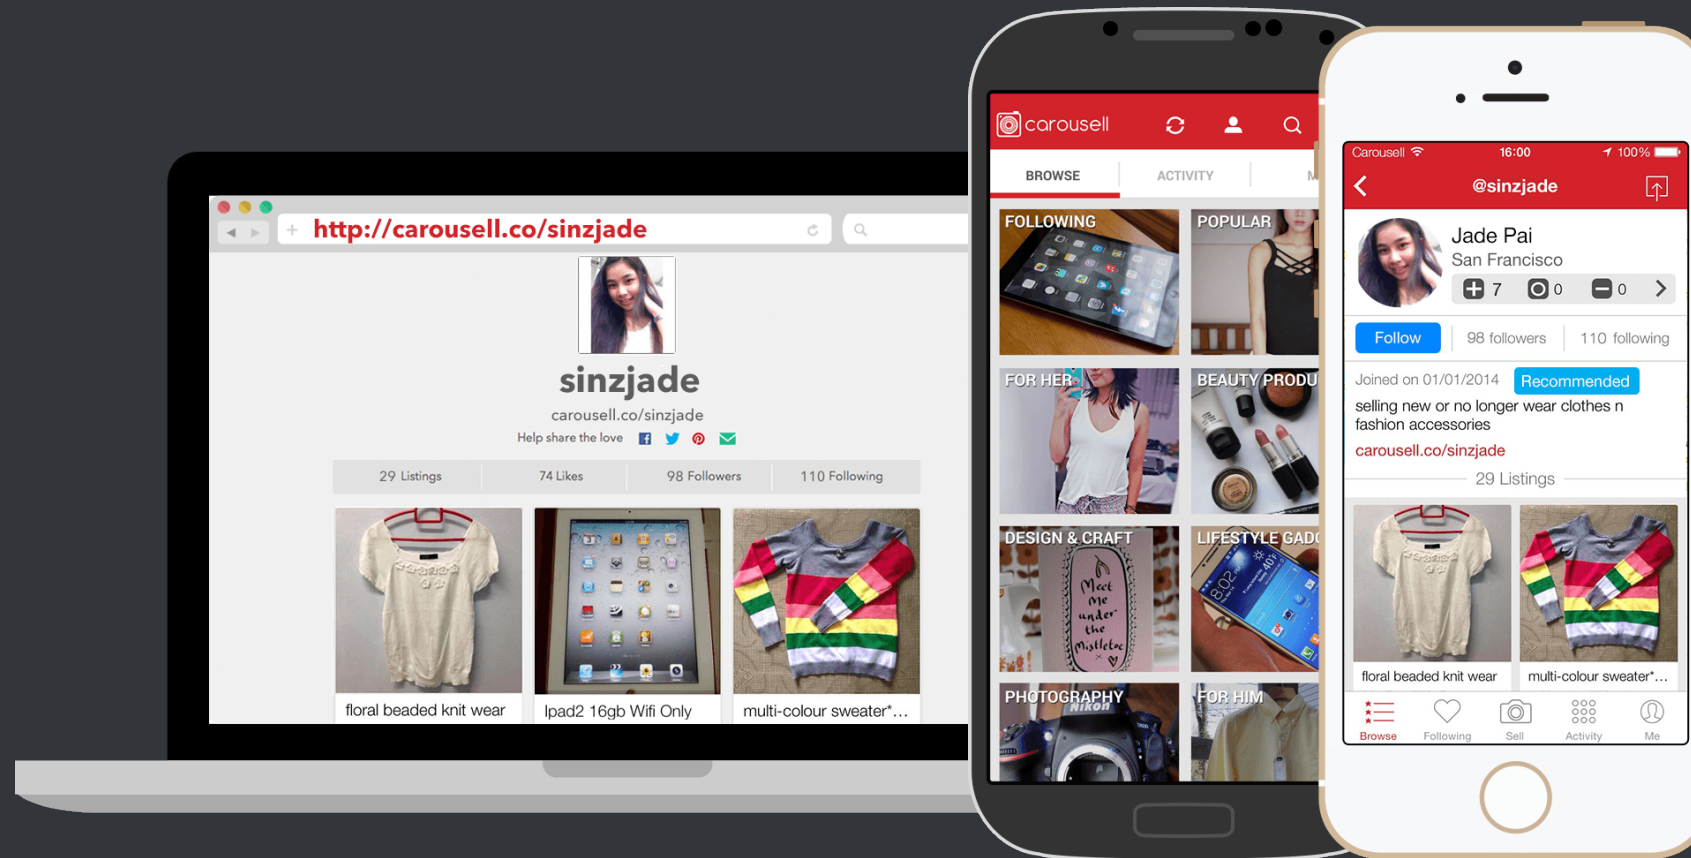
Tips, tricks and tuning
-beyond-
simple Django projects.

**premature optimization
is the root of all evil**

**yet we should not pass up our
opportunities in that critical 3%**

hi there, I am
victor neo

Lead Engineer @ Carousell



→ Caching *
ORM / Database *
Asynchronous Tasks
Logging / Monitoring

LWN featured content

[\$] What's in a (CentOS) version number?

[Front] Posted Jun 11, 2014 14:45 UTC (Wed) by corbet

The CentOS project has made its reputation by doing one thing very well: repackaging the Red Hat Enterprise Linux (RHEL) distribution into a freely distributable form. For users who are able to do without the support services offered by Red Hat, CentOS has been an invaluable resource. It is perhaps not surprising that CentOS users worry about the future of this distribution; they are getting a lot for free and many of them know that such situations are not always sustainable. For CentOS, keeping its user base depends on maintaining a certain level of trust so that users know it will continue to be available, stable, and free. The discussion around a proposal on version numbers shows just how easy that trust could be to lose.

[Full Story](#) (comments: 45)

PGCon 2014: Clustering and VODKA

[Development] Posted Jun 4, 2014 18:49 UTC (Wed) by jake

The eighth annual PostgreSQL developer conference, known as [PGCon](#), concluded on May 24th in Ottawa, Canada. This event has stretched into five days of meetings, talks, and discussions for 230 members of the PostgreSQL core community, which consists both of contributors and database administrators. PGCon serves to focus the whole PostgreSQL development community on deciding what's going to be in next year's PostgreSQL release as well as on showing off new features that contributors have developed. This year's conference included meetings of the main PostgreSQL team as well as for the Postgres-XC team, a keynote by Dr. Richard Hipp, and new code to put VODKA in your database.

Subscribers can click below for the full report from guest author Josh Berkus.

[Full Story](#) (comments: 17)

Posts + Comments Count

Homepage of a Blog

1. GET /index
2. Load posts from DB
3. Render template

Takes 200~500ms :(

**Caching is your
first layer of defence**

Speeding it up

1. **GET /index**
2. **Return cached output**

(completes in < 50ms!)

Django's per view cache

```
@cache_page(60 * 15)  
def homepage(request):
```

Caches homepage for ~15 minutes

Memcached

By far the fastest, most efficient type of cache available to Django, **Memcached** is an entirely memory-based cache framework originally developed to handle high loads at LiveJournal.com and subsequently open-sourced by Danga Interactive. It is used by sites such as Facebook and Wikipedia to reduce database access and dramatically increase site performance.

But I need more
CONTROL!

Django's low level cache

```
posts = Post.objects....all()  
  
# some complicated logic here  
  
cache.set('posts', posts, 60)
```


Django Cache Machine

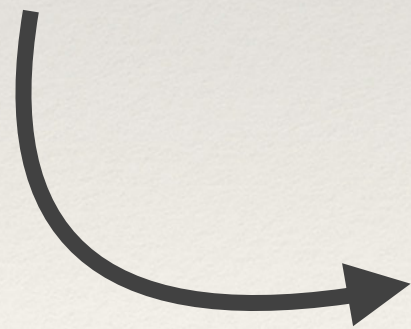
Automatic caching and invalidation of Models

Provides a Cache Manager

Cache Machine

```
# models.py  
class Post(models.Model):  
    objects = CachingManager()
```

```
Post.objects....all()
```



**Subsequent calls
are cached!**

Can I use
REDIS?

Redis for Caching

Django-Redis: Use redis instead of Memcached

**Django-cacheops: Similar to Cache Machine,
for Redis**

Decisions, decisions

	Memcached	Redis
Feature-packed	Cache Machine	Django-Cacheops
Custom	Roll your own with Django cache API	Django-Redis

Rule of thumb

Cache if data freshness is not an issue

Critical if computation is expensive

Caching

→ **ORM / Database**

Asynchronous Tasks

Logging / Monitoring

Foreign Keys

```
e = Post.objects.get(id=5)
```

```
e.blog.name      # Additional  
                  # DB query  
                  # for blog
```


`select_related`

```
e = Entry.objects.\n    select_related('blog').\n    get(id=5)
```

```
e.blog.name      # No DB query!
```


select vs prefetch related

**One-to-one / Foreign Key:
select_related**

**Many-to-many / Many-to-one / Generic relations:
prefetch_related**


```
Post.objects.filter(  
    is_published=True,  
    is_edited=True,  
...).all()
```



```
SELECT ... FROM blog_posts  
WHERE is_published = true  
AND is_edited = true  
AND ...
```


**Know thy
Database !**

PostgreSQL

EXPLAIN ANALYZE <query>

Non-Active Users

```
User.objects.\nfilter(is_active=False).\n.all()
```


Non-Active Users (SQL)

```
SELECT * FROM auth_users  
WHERE is_active = false;
```


Analyzing Queries

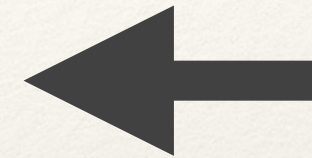


Prepend

```
EXPLAIN ANALYZE  
SELECT * FROM AUTH_USERS  
WHERE is_active = false;
```


QUERY PLAN

Seq Scan on auth_user



(cost=0.00..15761.01 rows=2413 width=140)

(actual time=0.161..279.318 rows=2384 loops=1)

Filter: (NOT is_active)

Total runtime: 280.890 ms



(3 rows)

Create an Index on is_active attribute



```
CREATE INDEX CONCURRENTLY  
ON auth_user (is_active)  
WHERE is_active = false;
```


QUERY PLAN

**Index Scan using
auth_user_is_active_idx on ←
auth_user**

(cost=0.00..59.19 rows=2413 width=140)
(actual time=0.129..8.824 rows=2384 loops=1)

Index Cond: (is_active = false)

Total runtime: 9.779 ms ←

(3 rows)

Postgres Weekly

A free, once-weekly e-mail round-up of PostgreSQL news and articles



Enter your e-mail address

Sign Me Up!

**Wait, how do I view
the SQL queries?**

Django debug toolbar

Select user to change | Django site admin

Select user to change | Django ...

SQL queries from 1 connection

default

5.02 ms (5 queries)

Query	Timeline	Time (ms)	Action
<div><div></div><div><div>SELECT ... FROM "django_session"</div><div>WHERE</div><div>("django_session"."session_key" = 'n30ikmlcaz0e7j1hte50a4x01z5hwt6c' AND "django_session"."expire_date" > '2013-12-21 03:26:09.232862')</div></div></div>	<div></div>	1.71	<div>Sel</div> <div>Expl</div>
<div><div></div><div><div>SELECT ... FROM "auth_user"</div><div>WHERE "auth_user"."id" = 1</div></div></div>	<div></div>	1.01	<div>Sel</div> <div>Expl</div>
<div><div></div><div><div>SELECT ... FROM "auth_group"</div></div></div>	<div></div>	0.43	<div>Sel</div> <div>Expl</div>
<div><div></div><div><div>SELECT ... FROM "auth_user"</div></div></div>	<div></div>	0.68	<div>Sel</div> <div>Expl</div>
<div><div></div><div><div>SELECT ... FROM "auth_user"</div><div>ORDER BY "auth_user"."username"</div><div>ASC, "auth_user"."id" DESC</div></div></div>	<div></div>	1.19	<div>Sel</div> <div>Expl</div>

Hide »

Versions

DJANGO 1.6

Time

Settings

Headers

Request

SQL

Templates

Static files

Silk

Summary

1081	2821	70ms	0.85	2ms
Requests	Profiles	Avg. Time	Avg. #Queries	Avg. DB Time

Most Time Overall

2014.06.16 10:39.424	2014.06.14 13:18.467	2014.06.15 02:19.287	2014.06.17 22:06.094	2014.06.15 04:38.660
200 GET /projects/silk/	200 GET /	200 GET /blog/2/	200 GET /rss	500 GET /blog/
743ms overall +6ms 0ms on queries +6ms 0 queries +1	696ms overall +486ms 13ms on queries +13ms 1 queries +5	580ms overall +82ms 19ms on queries +37ms 1 queries +5	93ms overall +8ms 2ms on queries +2ms 1 queries +2	30ms overall +2ms 0ms on queries +1ms 0 queries +1

Most Time Spent in Database

2014.06.16 17:52.780	2014.06.17 19:18.214	2014.06.15 04:29.764	2014.06.16 18:36.431	2014.06.17 11:46.400
200 GET /projects/silk/	500 GET /blog/	200 GET /blog/2/	403 POST /	200 GET /rss
17ms overall +2ms 0ms on queries +2ms 0 queries +1	16ms overall +1ms 0ms on queries +1ms 0 queries +1	252ms overall +42ms 24ms on queries +35ms 1 queries +5	21ms overall +1ms 0ms on queries +1ms 0 queries +1	48ms overall +16ms 3ms on queries +3ms 1 queries +2

Django Query Inspector

[SQL] repeated query (6x): SELECT "customer_role"."id",
"customer_role"."contact_id", "customer_role"."name"
FROM "customer_role" WHERE "customer_role"."contact_id" = ?

Suitable for API projects with no web UI

Search Engine?

PostgreSQL comes with full-text search

Heavy search traffic?


Consider Elasticsearch, Solr



Haystack

Modular search for django

Find the needle you're looking for.

 **Download**

Haystack

Supports Elasticsearch, Solr and more

Easy to get started with `manage.py` commands

Familiar ORM syntax for searching

Doing a search

```
SearchQuerySet().models(Post).  
    filter(content='Python').all()
```

```
[  
<SearchResult: blog.post (pk=u'1')>,  
<SearchResult: blog.post (pk=u'2')>,  
...]
```


Happy ORM

Use select/prefetch related to reduce queries

Understand your DB's query planner

Haystack for search

Caching

ORM / Database

➔ **Asynchronous Tasks**

Logging / Monitoring



Celery: Distributed Task Queue

Celery is an asynchronous task queue/job queue based on distributed message passing. It is focused on real-time operation, but supports scheduling as well.

The execution units, called tasks, are executed concurrently on a single or more worker servers using multiprocessing, Eventlet, or gevent. Tasks can execute asynchronously (in the background) or synchronously (wait until ready).

Celery is used in production systems to process millions of tasks a day.

View is slow :(

```
reset_pw_email(user.email)
```

```
# needs to wait for email to  
# be sent before we can send  
# a response
```

```
return HttpResponse(...)
```


Celery Tasks

```
@task  
def reset_pw_email(email):  
    ...
```

Just add @task decorator!

Calling Tasks

In your Django view

reset_pw_email.delay(user.email)

**View continues,
without waiting for email to be sent**

HTTP Callback tasks

```
HttpDispatchTask.delay(  
    url='http://a.com/multiply',  
    method='GET', x=10, y=10)
```

Awesome for microservices

Caching

ORM / Database

Asynchronous Tasks

➔ **Logging / Monitoring**

django.db.models.query in get
MultipleObjectsReturned: get()
returned more than one User --
it returned 2!

Errors Happen
Let's deal with them

Watch them as they happen



django.db.models.query in get

MultipleObjectsReturned: get() returned more than one User -- it returned 2!

users 1 2 hours ago root



Log errors when they happen in production

Open source Django project

Awesome web interface



Stream

Settings

Search query or event ID



Resolve Feed

Pause Updates

Sort by: Last Seen ▾

Between: The Past and The Present ▾



No events to show.

We'll notify you if that changes. In the meantime why not take a moment to become more familiar with Sentry.

[Installation instructions](#) [Project settings](#)



BOOKMARKS

All Events

Only Bookmarks

STATUS

Unresolved



OS

Select a os



FOO

Select a foo



VERSION

Select a version



LEVEL

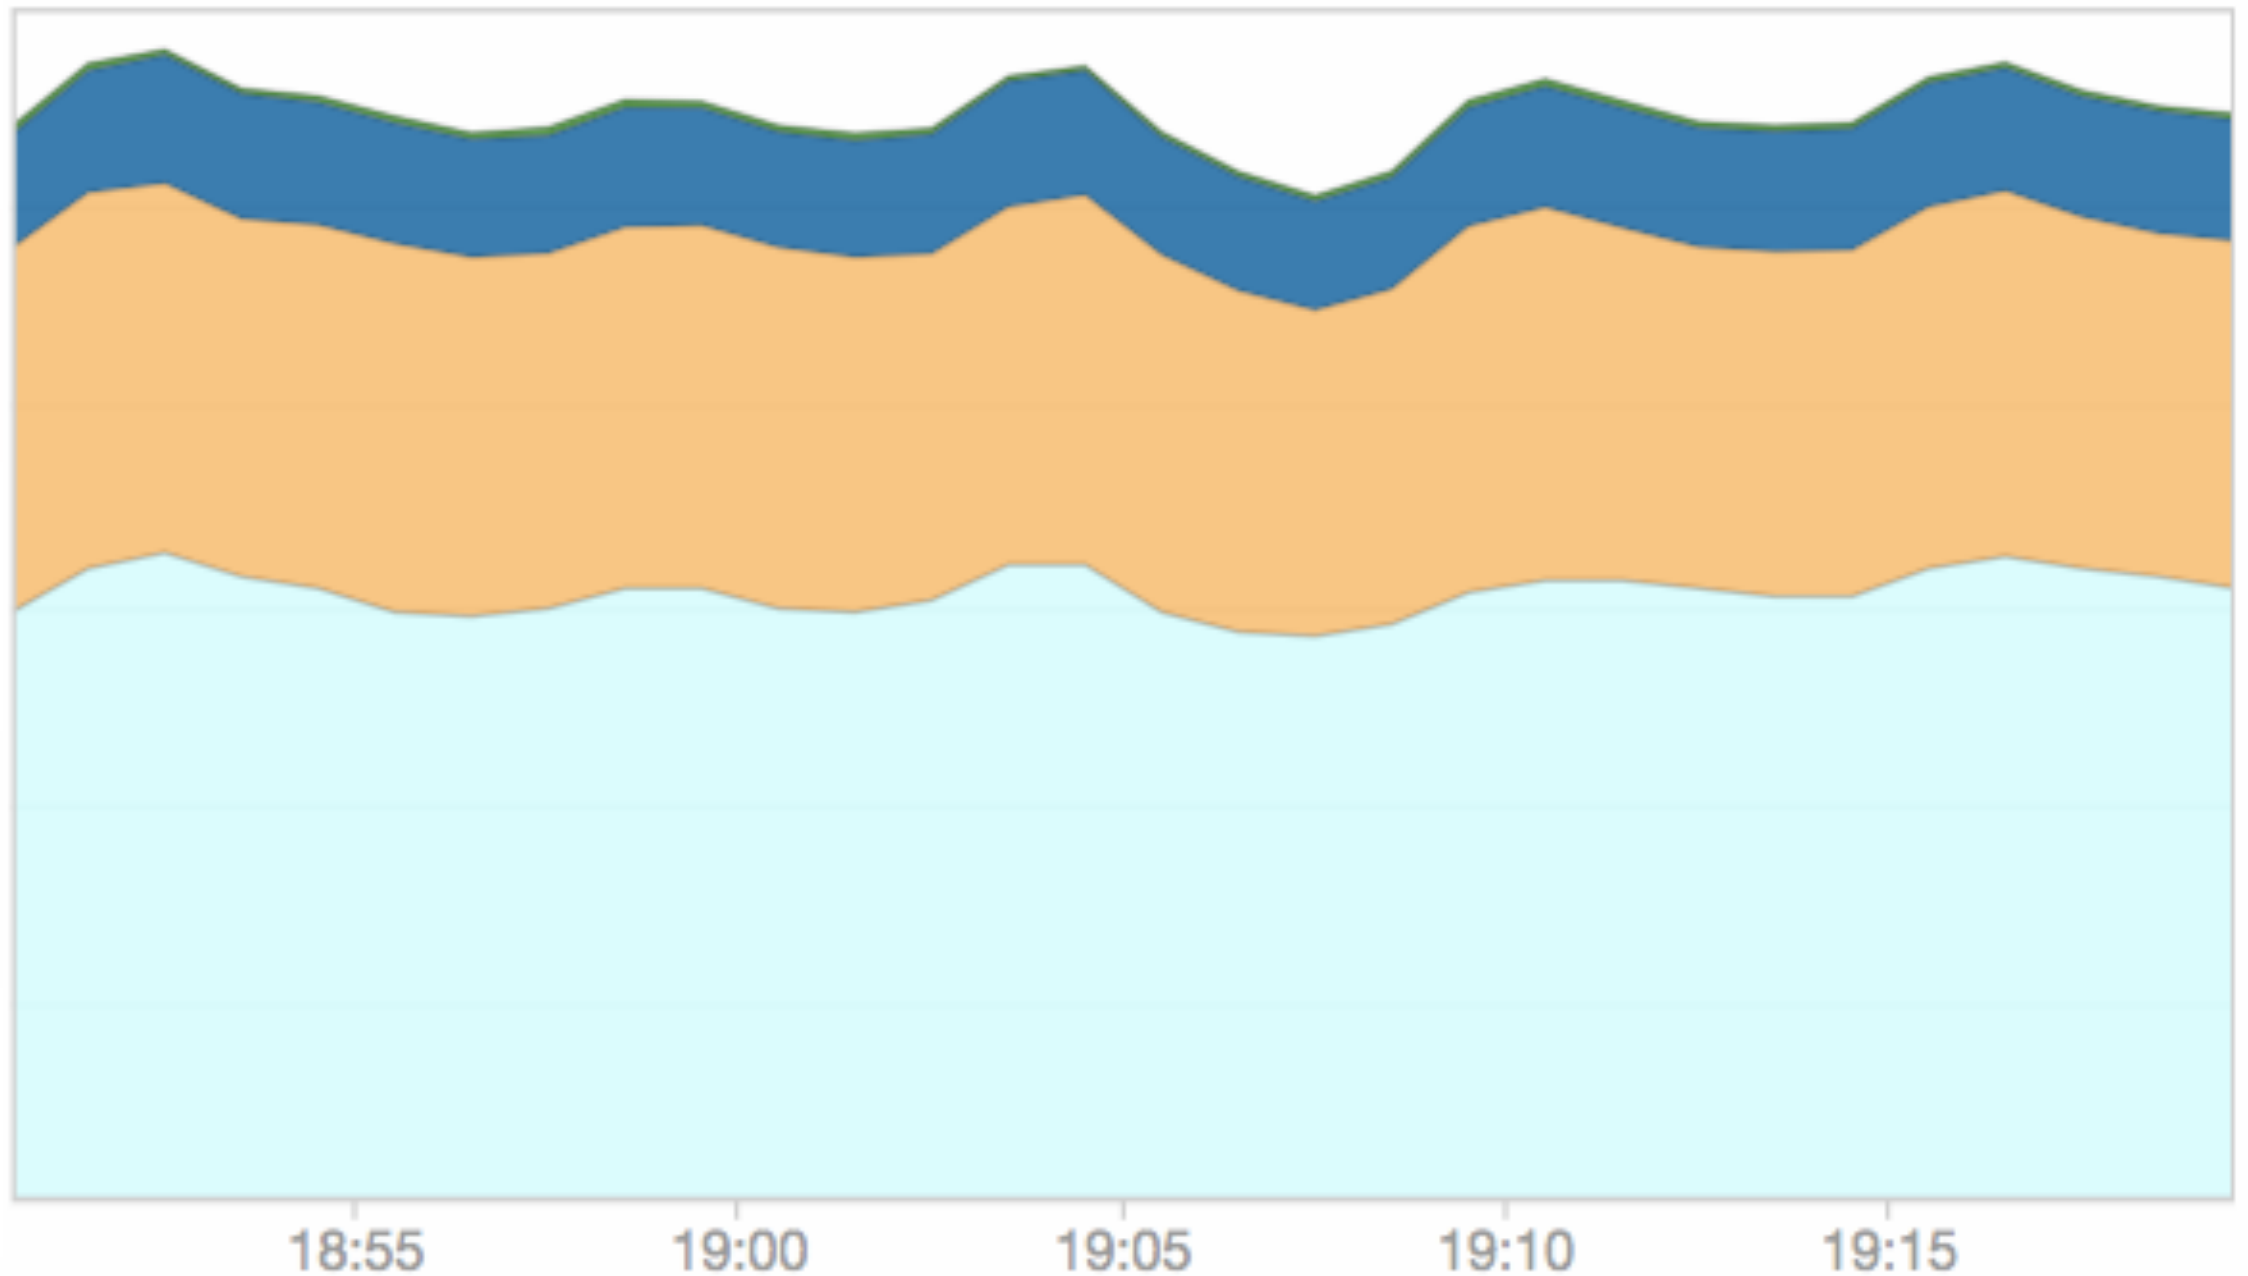
Select a level



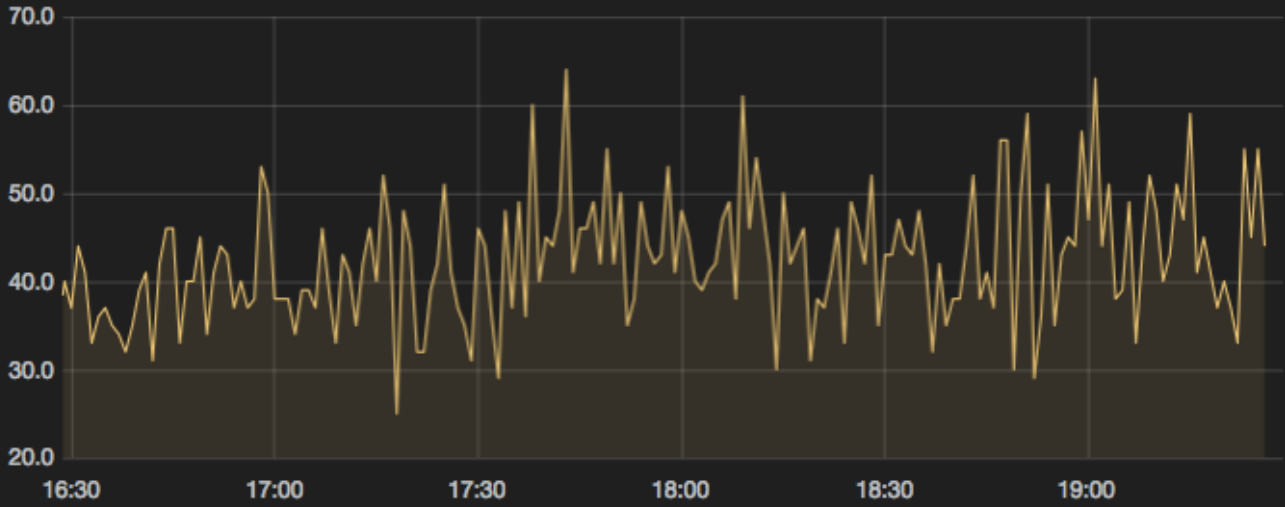
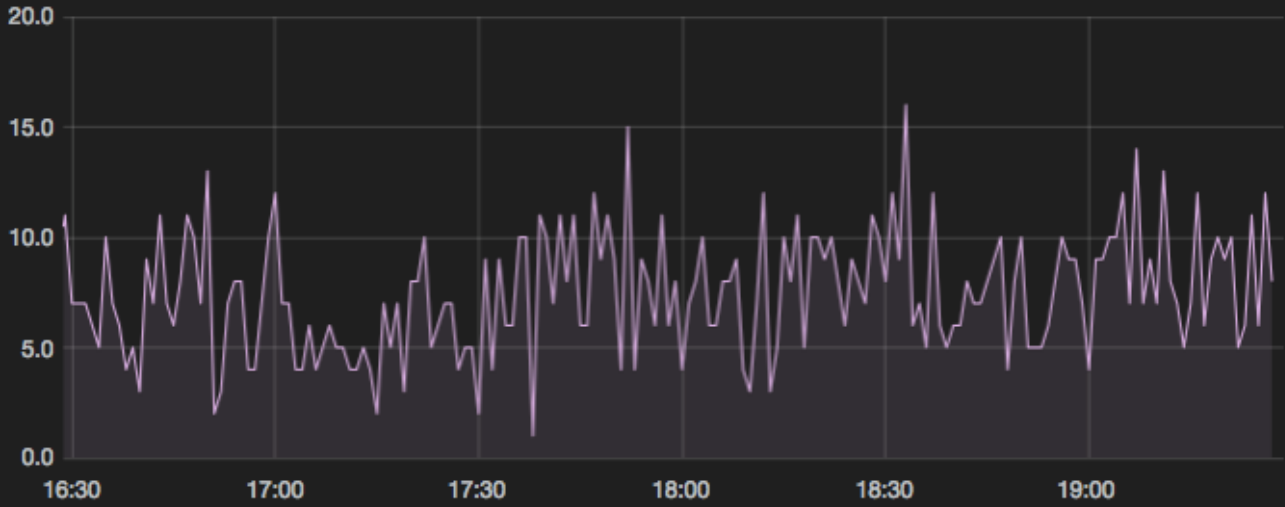
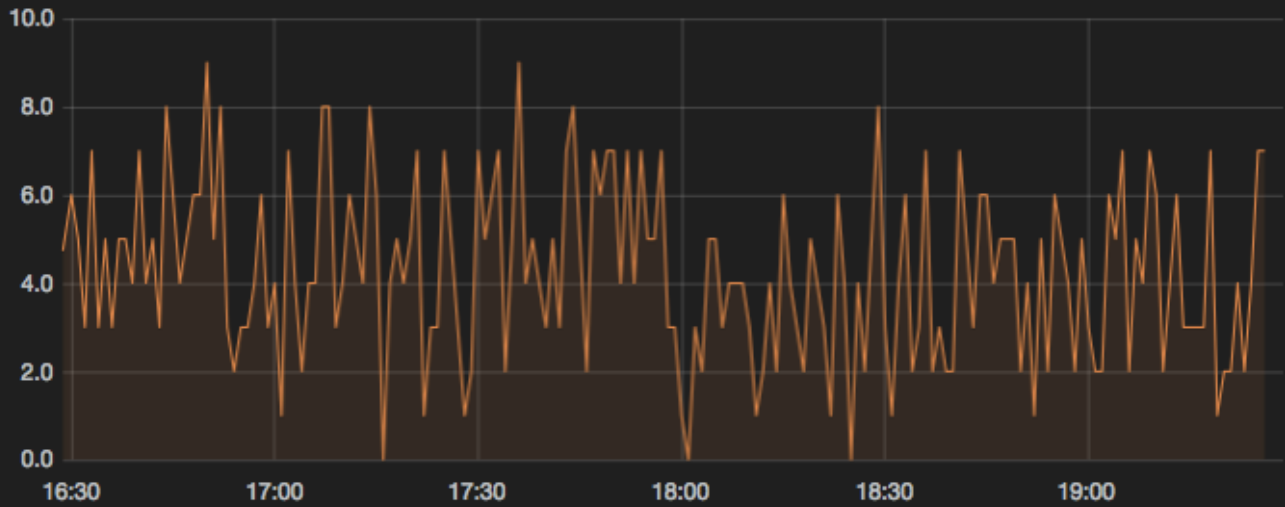
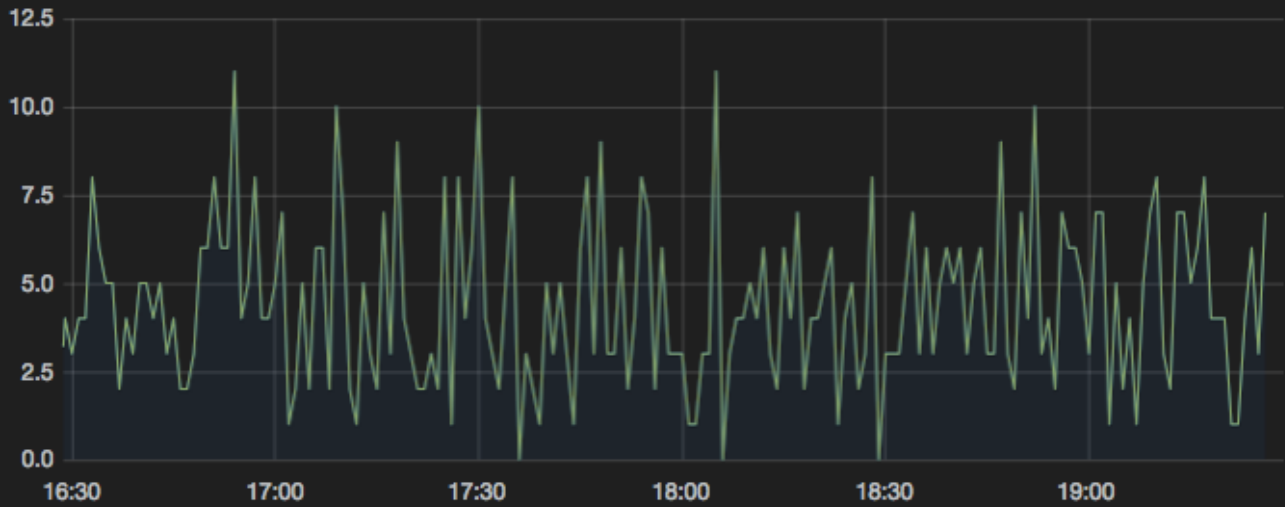
Monitoring
or “is someone using my app”

New Relic

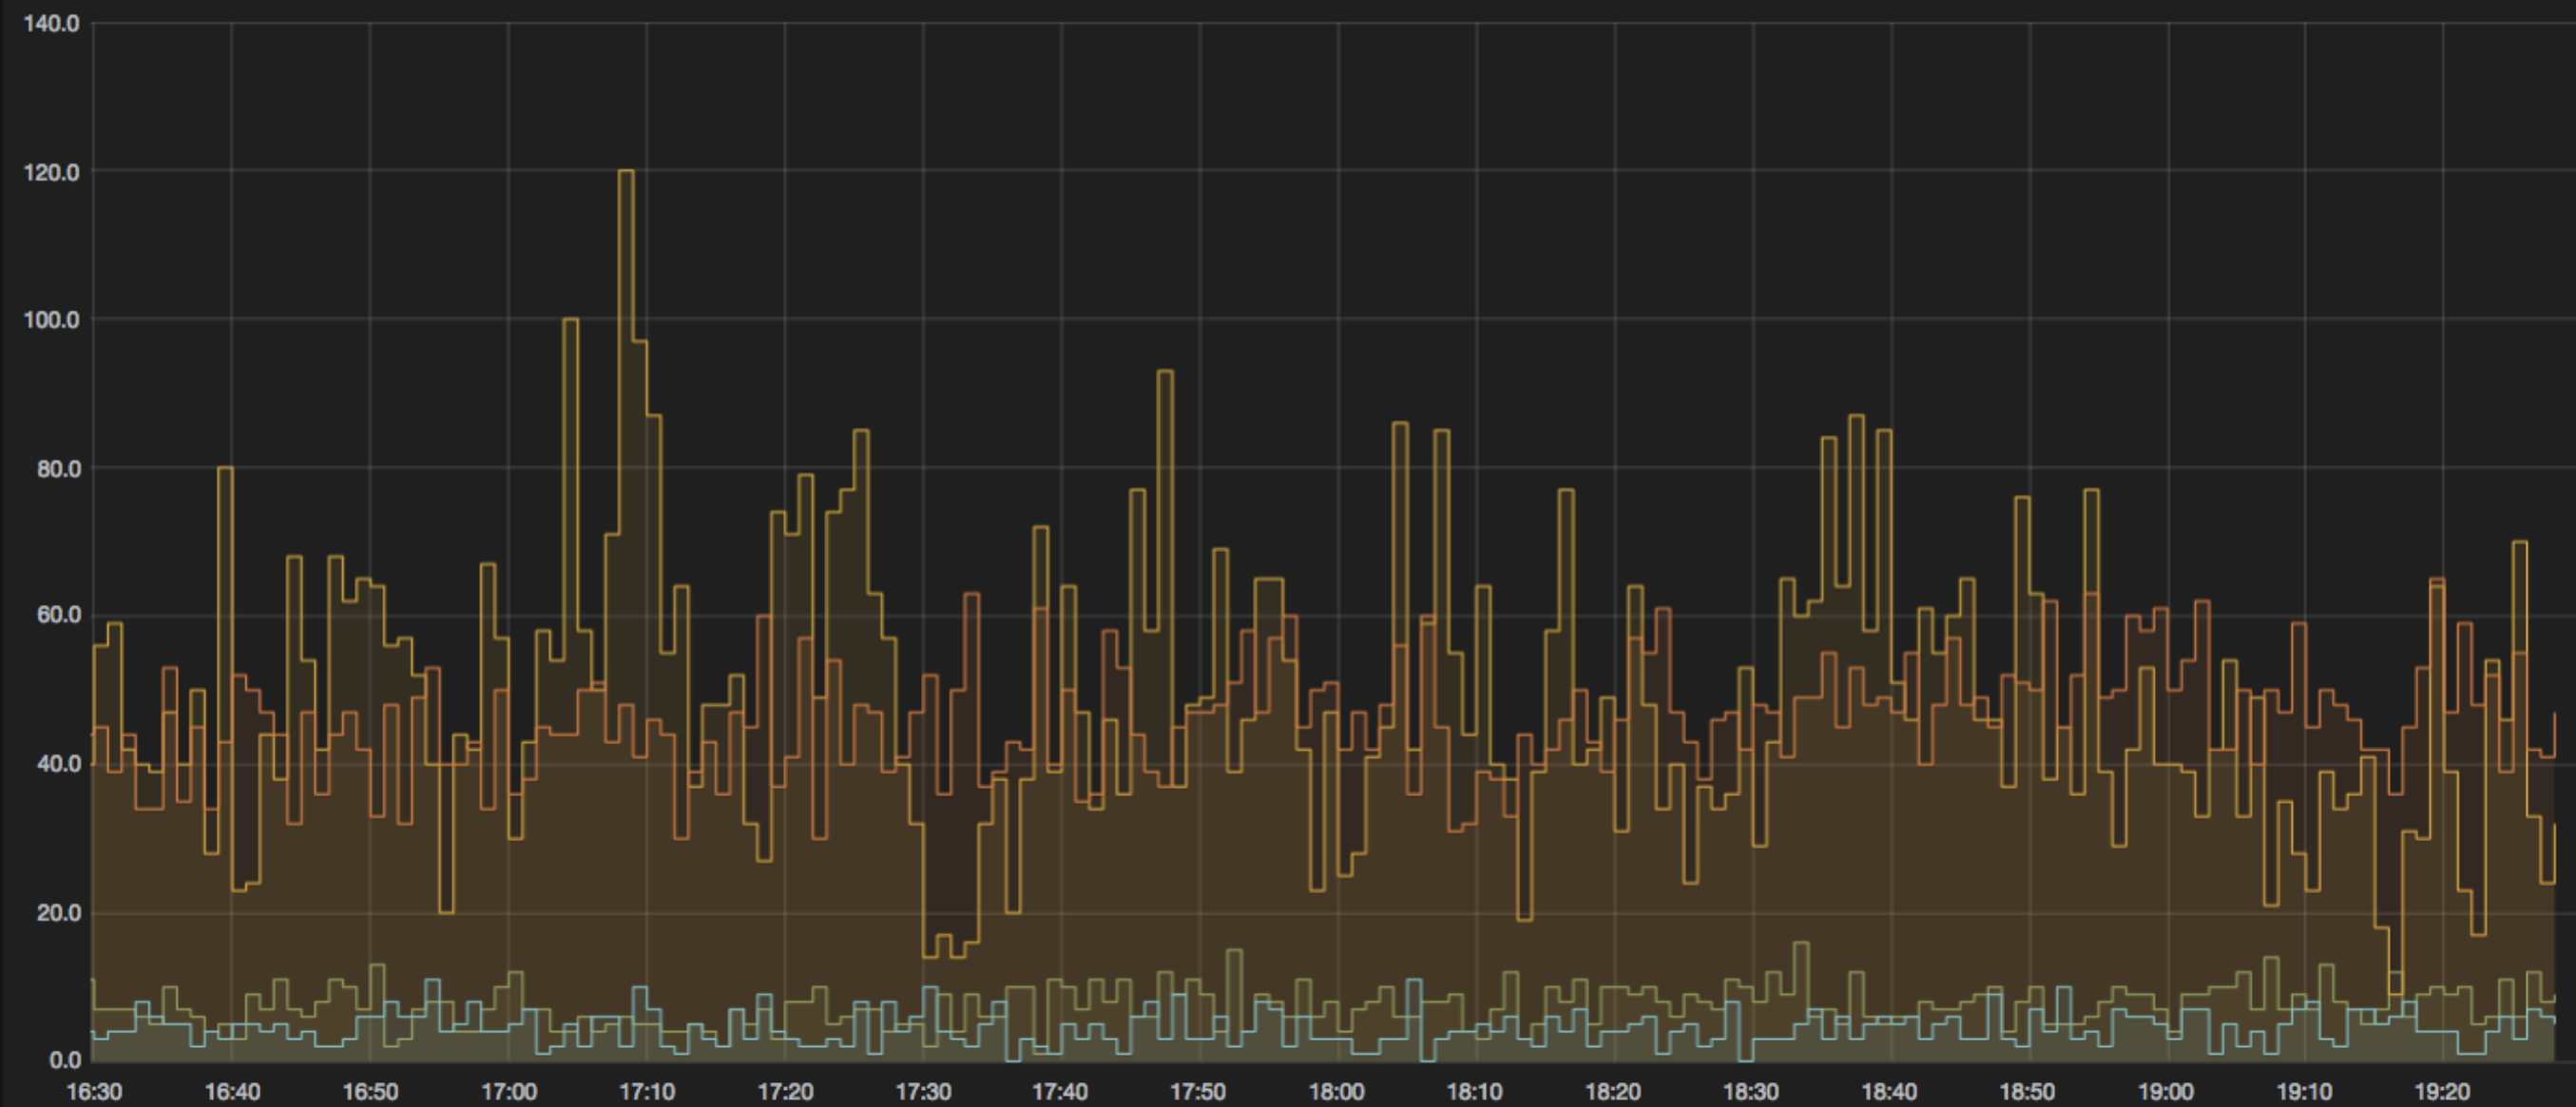
Web transactions response time ▾



Instrument Everything



Everything



Logging activity

```
c = statsd.Counter( 'UserSignups' )  
c.increment()
```

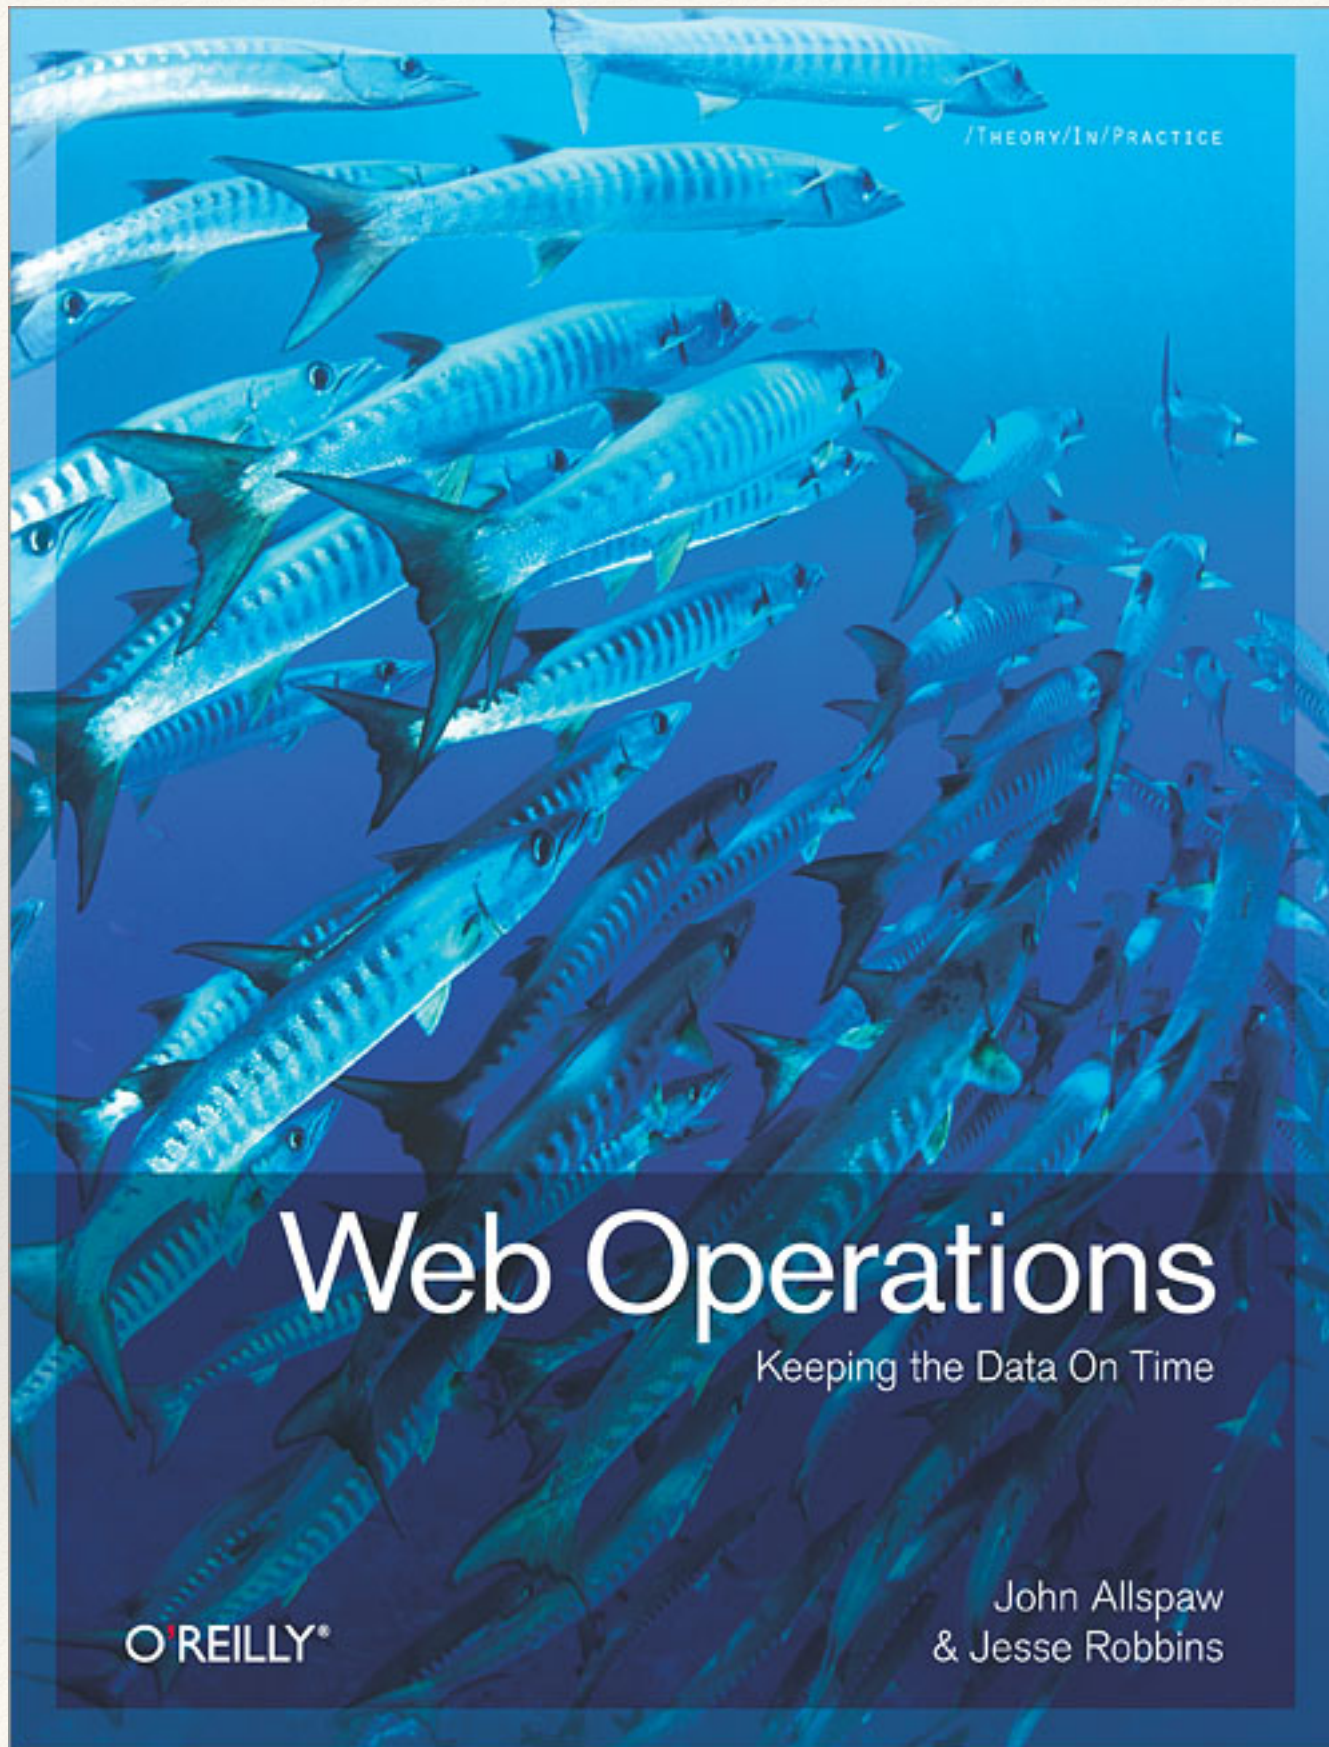

What to track?

Technical Metrics:

Cache misses, images uploaded

Business Metrics:

Number of signups



Web Operations

Keeping the Data On Time

Caching: memcached, redis

ORM / Database: FK keys, DB queries, Haystack

Asynchronous Tasks: Celery

Logging / Monitoring: Sentry, Graphite

Thank You!

come chat with me anytime

@victorneo